

1 **1. Introduction**

2 Date of Discovery Workshop: 2/24/26

3 Customer Name: FinServe Analytics

4

5 **2. PoC Scope**

6 **2.1. Background**

7 FinServe Analytics, a leading company in the financial services industry, recognized the need to transform how their
8 Relationship Managers (RMs) access and utilize critical insurance product information. The company manages an extensive
9 portfolio of insurance products, each documented in comprehensive materials that can span 700-1000 pages. These complex
10 documents contain detailed product specifications, terms and conditions, pricing structures, coverage details, and regulatory
11 compliance information.

12 Relationship Managers frequently need to retrieve specific information quickly during customer interactions to provide
13 accurate guidance and recommendations. The traditional approach of manually searching through lengthy documents was
14 time-consuming, prone to errors, and created friction in the customer experience. FinServe Analytics aimed to leverage
15 Generative AI technology to create an intelligent solution that would empower RMs with instant access to accurate product
16 information, ultimately enhancing customer service quality and operational efficiency.

17 The company operates in a highly regulated financial services environment where accuracy, compliance, and customer trust
18 are paramount. RMs serve as the primary interface between FinServe Analytics and its clients, requiring deep knowledge of
19 complex insurance products to provide personalized recommendations and answer detailed questions. The sheer volume and
20 complexity of product documentation created significant operational challenges, with RMs spending considerable time
21 searching through documents, leading to delayed responses, potential inaccuracies, and reduced time available for high-value
22 customer interactions.

23

24 **2.2. POC Objectives**

25 The primary objective is to develop an Agentic AI-powered chatbot built on AWS services that can autonomously assist
26 Relationship Managers in accessing and synthesizing complex insurance product information. The solution leverages
27 Foundation Models and the Retrieval Augmented Generation (RAG) paradigm to create an intelligent agent capable of
28 understanding context, reasoning about information needs, and taking autonomous actions.

29 Key objectives include:

- 30
- 31 • Dramatically reduce information retrieval time for Relationship Managers
 - 32 • Ensure consistent accuracy in product information delivery
 - 33 • Enable RMs to serve customers more effectively with near-real time access to comprehensive product knowledge
 - 34 • Maintain strict compliance with financial services regulations

35

36 **2.3. POC Assumptions & Limitations**

37 Data Availability:

- 38
- 39 • Output data: Natural language responses in Italian containing synthesized insurance product information,
40 recommendations, and guidance for Relationship Managers
 - 41 • Input data: Comprehensive insurance product documentation spanning 700-1000 pages per product, including
42 product specifications, terms and conditions, pricing structures, coverage details, and regulatory compliance
43 information. Information is spread across multiple documents in various formats, including both structured data
44 (tables, specifications) and unstructured content (policy descriptions, terms and conditions).
 - 45 ○ Document types:
 - 46 ■ PDFs
 - 47 ■ Word Documents

- 55 Assumptions:
- 56 • The solution will operate in Italian language
- 57 • RMs will interact with the system during customer-facing interactions
- 58 • The system must handle both structured (tables, specifications) and unstructured data (policy descriptions, terms and
- 59 conditions)
- 60 • Multi-turn conversational capability is required
- 61 • Context retention across dialogue history is necessary
- 62
- 63 Model Evaluation [Critical to Success]:
- 64 • Qualitative Assessment
- 65 • Accuracy of information retrieval from 700-1000 page documents
- 66 • Relevance of responses to RM queries
- 67 • Quality of synthesized answers from multiple document sections
- 68 • Appropriateness of response complexity based on RM expertise level
- 69 • Compliance with financial services regulations
- 70 • Natural conversational flow in Italian
- 71
- 72 • Quantitative Assessment
- 73 • RM productivity improvements
- 74 • Customer satisfaction scores
- 75 • Response time reductions
- 76 • Return on investment for the AI implementation
- 77 • I don't know specific baseline or target metrics
- 78
- 79

80 2.4. POC Challenges and Opportunities

- 81
- 82 Challenges:
- 83 • Managing extremely large documents (700-1000 pages per product)
- 84 • Ensuring accuracy in a highly regulated financial services environment
- 85 • Synthesizing information spread across multiple documents in various formats
- 86 • Maintaining data security and regulatory compliance
- 87 • Handling complex insurance product terminology and concepts
- 88 • Providing responses in Italian with appropriate context
- 89 Opportunities:
- 90 • Transformative improvement in RM efficiency and customer service quality
- 91 • Competitive differentiation through advanced AI technology
- 92 • Reduction in errors and improved consistency
- 93 • Enhanced customer experience through faster, more accurate responses
- 94 • Foundation for broader AI adoption across the organization
- 95
- 96

97 2.5. POC Environment & Data Requirements

- 98
- 99 Environment Requirements:
- 100 • AWS services infrastructure
- 101 • Support for Foundation Models
- 102 • RAG (Retrieval Augmented Generation) implementation capability
- 103 • Multi-agent architecture support
- 104 • Italian language processing capability
- 105 • Integration with RM workflows
- 106 • Advanced data protection mechanisms
- 107 • Harmful content filtering

- 108 • Compliance monitoring capabilities

109

110 Data Requirements:

- 111 • Access to comprehensive insurance product documentation (700-1000 pages per product)
- 112 • Product specifications, terms and conditions, pricing structures, coverage details
- 113 • Regulatory compliance information
- 114 • Multiple document formats (structured and unstructured)

115

116 **2.6. POC Success Criteria, Evaluation Method and Metrics**

117

118 Business Metrics:

- 119 • RM productivity improvements
- 120 • Customer satisfaction scores
- 121 • Response time reductions
- 122 • Return on investment
- 123 • Reduction in information retrieval time
- 124 • Increase in time available for high-value customer interactions
- 125 • Reduction in errors and inaccuracies

126 Technical Metrics:

- 127 • Accuracy of information retrieval
- 128 • Relevance of responses
- 129 • Response time/latency
- 130 • System availability and reliability
- 131 • Context retention across multi-turn conversations

132 Operational Metrics:

- 133 • User adoption rate among RMs
- 134 • Number of queries handled
- 135 • Successful resolution rate
- 136 • Compliance adherence rate

137 Evaluation Method:

- 138 • Comprehensive evaluation framework to measure accuracy, relevance, response time, and user satisfaction
- 139 • Continuous evaluation and refinement based on feedback
- 140 • Learning from interaction patterns to improve performance over time

141

142

143 **2.7. Deliverables**

144 Minimum Lovable Product (MLP) includes:

- 145 • Conversational chatbot interface integrated into RM workflows
- 146 • Intelligent knowledge retrieval system powered by RAG
- 147 • Multi-turn conversation capability with context retention
- 148 • Comprehensive evaluation framework to measure accuracy, relevance, response time, and user satisfaction

149 Additional deliverables:

- 150 • Intelligent Agent Architecture for navigating insurance documentation
- 151 • Autonomous Conversational Experience with multi-turn conversations in Italian
- 152 • Multi-Agent Knowledge Retrieval System with specialized collaborative agents (document retrieval agents, extraction agents, validation agents, synthesis agents)
- 153
- 154 • Advanced data protection mechanisms and harmful content filtering

155

156

157

158

159

160 **2.8. Timeline**

161
162
163
164

2.9. Future Scope

165 The goal is to create an autonomous AI ecosystem that:
166 • Continuously improves through learning from interactions
167 • Scales across the organization for broader AI adoption
168 • Expands to additional use cases beyond insurance product support
169 • Integrates with additional systems and workflows
170 • Enhances capabilities based on user feedback and evolving needs
171 • Create a system that continuously improves through learning from interactions
172
173 • Dynamic user persona adaptation that adjusts response complexity and detail based on RM expertise
174
175
176

3. Project Execution

3.1. POC Execution

179
180

3.2. Proofpoint Business and Technical Owners

181 Executive Sponsor: Chief Innovation Officer

182 • Responsible for driving adoption of cutting-edge AI technologies across FinServe Analytics' customer-facing operations
183 • Ensures alignment with the company's digital transformation roadmap
184 • Oversees resource allocation and manages stakeholder expectations
185 • Works closely with Head of Relationship Management, IT leadership, and compliance teams
186 • Accountable for measuring success metrics
187

188

189 Subject Matter Experts: Relationship Managers and relevant product/compliance experts (specific individuals not identified)

190

191

3.3. Next Steps

192
193
194

4. Additional Notes and Assumptions

195 Additional Notes and Assumptions

196 • The solution must operate in Italian language
197 • Financial services regulatory compliance is paramount
198 • Data security and privacy protection are critical requirements
199 • The system must handle both structured and unstructured data
200 • Multi-agent collaborative architecture is planned
201 • Continuous learning and improvement capability is required
202 • Integration with existing RM workflows is essential
203
204
205
206

5. Appendix

5.1. Preliminary Resource Requirements

207
208
209

5.2. Discovery Workshop Participants

210
211
212
213

5.3. Customer Participants

215
216
217 **5.4. AWS Participants**
218
219
220