

Product Requirements Document (PRD)
Insurance Product Support – Minimum Lovable Product (MLP)

Date: February 24, 2026

Version: 0.1 (Draft)

1. Summary

This PRD defines the Minimum Lovable Product (MLP) described in the Insurance Product Support brief. The MLP is an initial deployment of a conversational chatbot for Relationship Managers (RMs), powered by Retrieval Augmented Generation (RAG), with multi-turn context retention, dynamic persona adaptation, and an evaluation framework measuring accuracy, relevance, response time, and user satisfaction.

2. Background and Problem

FinServe Analytics manages an extensive portfolio of insurance products, each documented in comprehensive materials spanning ~700–1000 pages per product. RMs must retrieve accurate, detailed information quickly during customer interactions. Manual search across long, multi-document materials is slow, error-prone, and degrades customer experience.

3. Goals

- Provide RMs with instant access to accurate insurance product information during customer interactions.
- Reduce time spent searching and synthesizing across long product documents.
- Support multi-turn RM conversations with context retention (Italian language).
- Adapt response complexity to RM expertise (persona adaptation).
- Establish an evaluation framework to measure accuracy, relevance, response time, and user satisfaction.

4. Non-Goals

- Replacing formal compliance or legal review processes.
- Fully autonomous multi-agent orchestration beyond the MLP scope (covered separately in the brief).
- Automated changes to product documentation or source systems.
- Customer-facing deployment (MLP is RM-facing).

5. Users and Personas

Primary user: Relationship Managers (RMs) serving clients in Italian. Personas vary by product expertise.

Persona	Typical Need	Response Style (Expected)
New RM	Quick, safe guidance on common questions	More explanatory; definitions; examples; high-level first
Experienced RM	Fast lookup of specific terms, limits, exclusions	Concise; structured; direct; assumes context
Specialist RM	Deep detail and edge cases; cross-reference sections	Detailed; cites sources; highlights conditions/exceptions

6. Primary Use Case

An RM asks a question about an insurance product (e.g., coverage, exclusions, pricing structures, terms and conditions, regulatory constraints). The chatbot retrieves the most relevant source passages across product documentation and returns a grounded answer in Italian, retaining conversational context across turns and adapting detail level to the RM persona.

7. Product Scope (MLP Components)

The MLP includes the following components as described in the brief:

- Conversational chatbot interface integrated into RM workflows
- Intelligent knowledge retrieval system powered by RAG
- Multi-turn conversation capability with context retention
- Dynamic user persona adaptation (response complexity/detail based on RM expertise)
- Evaluation framework measuring accuracy, relevance, response time, and user satisfaction

8. Functional Requirements

ID	Requirement	Description	Priority	Acceptance Criteria (MLP)
FR-01	Chat UI	RM can ask questions and receive answers in Italian; supports	P0	RM can submit query; receive answer in Italian within target

		copy/export of answers.		latency; basic conversation history visible.
FR-02	RAG Retrieval	System retrieves relevant passages from product documentation and uses them to ground responses.	P0	Answers include citations/snippets to retrieved passages; retrieval improves over naive keyword search baseline.
FR-03	Multi-turn Context	System retains context across multiple turns in the same conversation.	P0	Follow-up questions resolve pronouns/references correctly using conversation context; user can reset context.
FR-04	Persona Adaptation	System adjusts response complexity/detail based on RM persona/experience level.	P0	User can select persona or system infers from interaction; output changes in depth/format accordingly.
FR-05	Safety Controls	Data protection mechanisms and harmful content filtering consistent with regulated environment.	P0	System blocks/flags disallowed content; logs safety events; provides safe refusal messaging.
FR-06	Evaluation Framework	Measure accuracy, relevance, response time, and user satisfaction.	P0	Dashboards/reports exist for these metrics; periodic sampled review workflow defined.
FR-07	Workflow Integration (Lightweight)	Accessible from RM workflow entry point (e.g., intranet/portal/CRM	P1	Single sign-on or equivalent access; link from primary RM workspace;

link).

minimal friction to start a session.

9. Data and Content Requirements

Source information includes detailed product documentation spanning 700–1000 pages per product, often spread across multiple documents and formats. The solution must handle both structured information (tables/specifications) and unstructured content (policy descriptions, terms and conditions).

- Document ingestion and indexing suitable for RAG (chunking, metadata, versioning).
- Support for extracting/using structured tables where needed.
- Clear linkage between answers and source passages for auditability.

10. Compliance, Security, and Governance

- Operate in a highly regulated financial services context; accuracy, compliance, and customer trust are paramount.
- Implement data protection mechanisms and harmful content filtering.
- Maintain traceability from answer to source content to support review and audit needs.

11. Metrics and Measurement

MLP evaluation metrics (explicitly called out in the brief):

Metric	Definition	How Measured	Notes
Accuracy	Correctness of answers vs. source documents	Human review on sampled Q&A set; error categorization	Must account for compliance-critical errors
Relevance	Retrieved and presented info matches the question intent	Human rating + retrieval relevance scoring	Track top-k retrieval quality
Response Time	Latency from question submit to answer	System telemetry	Segment by doc size and complexity
User Satisfaction	RM perceived usefulness and trust	In-product feedback + periodic survey	Track adoption alongside satisfaction

Business success metrics referenced for executive sponsorship (tracked alongside MLP evaluation):

- RM productivity improvements
- Customer satisfaction scores
- Response time reductions
- Return on investment (ROI)

12. Release Plan and Milestones (Proposed)

Note: The source brief does not include dates. The milestones below are a proposed sequencing to deliver the MLP scope.

Milestone	Name	Exit Criteria
M1	Discovery & KPI definition	Confirm RM workflows, top question categories, and evaluation plan.
M2	Data ingestion & indexing	Ingest initial product docs; create RAG index; define doc versioning approach.
M3	MLP build	Chat UI, RAG, multi-turn context, persona adaptation, safety controls.
M4	Evaluation & pilot	Run evaluation framework with RM pilot; iterate on retrieval and responses.
M5	Production readiness	Monitoring, governance, runbook, training material, broader RM rollout.

13. Risks and Open Questions

- Document quality/format variability may degrade retrieval and grounding.
- Regulatory/compliance requirements may necessitate stricter refusal behavior, reducing perceived usefulness.
- Persona adaptation must be predictable; overly dynamic behavior may reduce trust.
- Evaluation coverage must include edge cases (exclusions, conditional clauses) to avoid high-impact mistakes.

Open questions (not answered in the brief):

- Which specific document formats are in scope (PDF, DOCX, scanned images, etc.)?
- What is the target RM population and rollout approach?
- What systems constitute “RM workflows” for integration (CRM, portal, intranet, etc.)?
- What are the compliance review and escalation requirements for answers?
- How frequently do product documents change, and what is the expected update SLA?